

VAL Vendor Analysis Assistant

Windows local deployment guide

This guide runs the current Streamlit application locally on Windows. The application uses the Azure OpenAI Python client with an Azure OpenAI endpoint, deployment name, API key, and API version. It does not require Azure CLI authentication or Azure AI Foundry thread access.

1. Prerequisites

Install:

- Python 3.11 or newer from <https://www.python.org/downloads/windows/>
- Git for Windows from <https://git-scm.com/download/win> (only if you need to clone the repository)

During Python installation, select Add Python to PATH.

2. Get the application

Open PowerShell and either clone the repository or change to the project folder:

```
git clone <repository-url>
cd VALAgent_Foundry
```

If the application work is on a feature branch, check out that branch before installing dependencies:

```
git checkout cursor/vendor-analysis-chat-a1f3
```

3. Create and activate a virtual environment

```
py -3 -m venv .venv
.\.venv\Scripts\Activate.ps1
python -m pip install --upgrade pip
pip install -r requirements.txt
```

If PowerShell blocks the activation script, allow it only for the current terminal session:

```
Set-ExecutionPolicy -Scope Process -ExecutionPolicy Bypass
.\.venv\Scripts\Activate.ps1
```

4. Configure Azure OpenAI credentials

Create a local configuration file:

```
Copy-Item .env.example .env
notepad .env
```

Replace every placeholder in `.env`:

```
AZURE_OPENAI_ENDPOINT=https://your-resource.openai.azure.com/
AZURE_OPENAI_DEPLOYMENT=your-deployment-name
AZURE_OPENAI_API_KEY=your-api-key
AZURE_OPENAI_API_VERSION=2024-12-01-preview
```

Use the endpoint and one of the API keys shown on the Azure OpenAI resource's Keys and Endpoint page. Use the deployment name from the Azure AI Foundry/Azure OpenAI deployment page; it can differ from the underlying model name.

Keep `.env` private. It is ignored by Git and must never be committed or shared. If an API key was exposed in source control, rotate it in Azure immediately and update `.env` with the replacement.

5. Run the application

With the virtual environment still active:

```
streamlit run app.py
```

Open the URL Streamlit prints, normally:

```
http://localhost:8501
```

Stop the server with `Ctrl+C`.

6. Verify the application

1. The sidebar should show Azure OpenAI configuration detected.
2. Confirm the sidebar model equals your deployment name.
3. Click New Conversation.
4. Send a prompt such as `Summarize the highest contract risks.`
5. Verify the assistant responds and that a follow-up question retains the prior chat context.

Troubleshooting

Symptom	Cause and resolution
`py` is not recognized	Reinstall Python and select **Add Python to PATH** , then open a new PowerShell window.
`streamlit` is not recognized	Activate `.env` and run `pip install -r requirements.txt` again.
Configuration required	Check all four `AZURE_OPENAI_*` values in `.env`; restart Streamlit after editing the file.
Azure OpenAI rejected the API key	Verify the key belongs to the resource in `AZURE_OPENAI_ENDPOINT`; rotate and replace the key if needed.
Deployment not found	Set `AZURE_OPENAI_DEPLOYMENT` to the Azure deployment name, not only the model name.
Network or firewall error	Configure Azure OpenAI resource networking to allow the Windows machine's network, or use the organization-approved private endpoint/VPN path.

Security checklist

- Store secrets only in `.env` or a managed secret store.
- Do not place credentials in `app.py`, `.env.example`, screenshots, tickets, or commits.
- Use a least-privilege Azure OpenAI key and rotate it according to your organization's policy.
- Restrict Azure OpenAI network access according to your organization's security requirements.